Research Article

Research and Politics
July-September 2020: 1–10

DOI: 10.1177/2053168020935250
journals.sagepub.com/home/rap
SAGE

Camouflaged propaganda: A survey experiment on political native advertising

Yaoyao Dai[1] and Luwei Luqiu[2]

Abstract

We examine a new form of propaganda, *political native advertising*, in which political actors, including foreign governments, buy space in independent media outlets to publish advertisements that are camouflaged as standard news stories. Those who engage in this form of propaganda hope to exploit the higher credibility of the hosting media site to enhance the persuasiveness of their message. Despite the obvious political implications and ethical issues at stake, political native advertising has received almost no scholarly attention. Our article begins to redress this imbalance. Using an online survey experiment with real political native advertisements in the *Washington Post* and *The Telegraph* bought by the Chinese government, we provide some of the first empirical evidence on basic but important features of political native advertising. We find, among other things, that respondents struggle to distinguish political advertisements from standard news stories regardless of their level of education and media literacy, that political advertisements are more convincing if they appear on and are perceived as news from an independent hosting media site than in a government-controlled news outlet, and that trust in the hosting media site declines if the political advertisement is detected.

Keywords

political native advertising, native advertising, China Watch, propaganda, disinformation campaign

Introduction

Since 2010, readers in the United States have been able to obtain news on China from a multi-page special section named *China Watch* in the *Washington Post*, the *Wall Street Journal*, and the *New York Times* (Cook, 2017; Fallows, 2011). Unfortunately, instead of being a special editorial column on China, the *China Watch* section is a paid supplement provided by *China Daily*, a Chinese government-controlled English-language newspaper (Fallows, 2011). As of March 2018, *China Daily* had cooperated with, and provided *China Watch* content to, more than 40 legacy news media in over 20 countries with a circulation of 4 million people.[1] This is all part of *China Daily*'s strategy to use the platforms and reputations of partnership publishers to increase the worldwide audience for its news stories (*China Daily*, 2018). China is not the only country that pays Western legacy media outlets to publish news stories from government-controlled media. For example, *Russia Beyond*, a Russian government-controlled media outlet, has also paid to place news stories in the *Washington Post* under the title *Russia Now*.[2] Political parties in democracies have also engaged in similarly deceptive advertising activities. In the United States, for example, both Democratic and Republican candidates have paid to insert campaign advertisements that mimic news stories and other forms of standard editorial content in domestic media outlets (Iversen and Knudsen, 2017; Dykhne, 2017; Murtha and Gourarie, 2016).

Communication and journalism scholars refer to paid content and advertisements camouflaged as standard editorial content as *native advertising* (Howe and Teufel, 2014). Unlike conventional sponsored content or advertisements, *native advertisements* are camouflaged as standard editorial content coming from the hosting media outlet. As a result, people are often unaware that they are reading sponsored and paid content. To date, existing studies have focused on native advertising almost exclusively in the context of commercial products (Carlson, 2015; Iversen and Knudsen, 2017; Jamieson et al., 2000; Batsell, 2017;

[1]Social Science Division, New York University Abu Dhabi, Abu Dhabi, United Arab Emirates
[2]School of Communication, Hong Kong Baptist University, Kowloon, HK, Hong Kong

Corresponding author:
Yaoyao Dai, Social Science Division, New York University Abu Dhabi, Saadiyat Island, Abu Dhabi, United Arab Emirates.
Email: yaoyao.dai@nyu.edu

Edmonds, 2017; Einstein, 2016; Mullin, 2017). We use the term *political native advertising* to refer to situations where political actors engage in native advertising.

Despite the obvious political implications and ethical concerns, such as co-opted (foreign) political influence and the threat to media freedom, political native advertising has received almost no scholarly attention. While there is a large and growing literature dealing with political propaganda, it tends to focus on issues of media control and the types of hard and heavy-handed propaganda that are easily detected, such as pro-regime reporting in government-controlled news programs and government-sponsored nationalist commercials (Di Tella and Franceschelli, 2011; Stockmann, 2013; Gehlbach and Sonin, 2014; Archer and Clinton, 2018; Huang, 2015, 2018; Little, 2017). A common claim in the literature is that propaganda, perhaps because it is often so easily detected, is not designed to achieve persuasion but instead to identify loyalists or signal the coercive capacity of the state (Crabtree et al., 2018; Little, 2018; Marquez, 2018; Huang, 2018, 2015). Relatively few studies have examined the type of soft propaganda that is harder to detect and is the hallmark of political native advertising. To our knowledge, there has been only one systematic study of political native advertising, which finds that clearly labeled hypothetical political native advertisements by political parties in Norway reduce citizen trust in political news in general (Iversen and Knudsen, 2017).

In this paper, we begin to redress this imbalance in the literature by providing one of the first investigations of political native advertising, particularly as it relates to its use by foreign governments. We use an online survey experiment with real political advertisements in the *Washington Post* and *The Telegraph* that were paid for by the Chinese government to obtain some basic and important, but as yet unknown, information about political native advertising. To what extent can readers detect political native advertising? Is political native advertising considered persuasive? Do hosting media outlets suffer a cost if the political native advertising is detected? We find that people often struggle to distinguish paid political advertisements from standard news stories regardless of their level of education and media literacy. We also find that the content found in political native advertisements is perceived to be more persuasive than the same content on a government-controlled media source when the readers perceive the political advertisements as editorial content from the hosting media outlet. Finally, we also find that people significantly reduce their level of trust in the hosting media outlet if they detect the political native advertising.

An experiment on political native advertising: *China Watch*

Given the lack of basic empirical information about political native advertising, we use an online survey experiment to examine the extent to which readers are able to detect political native advertisements, whether political native advertisements are considered persuasive, and whether independent hosting media outlets suffer a cost in terms of trust if readers detect a political native advertisement. These three basic elements are important in understanding political native advertising: the persuasiveness of political native advertising speaks to the motivation of the political sponsors; the reputational cost of accepting political native advertising is likely to influence the decision of independent media outlets to host such advertisements; the likelihood of being detected is likely to influence the persuasiveness of the political native advertisement and the actual reputational costs to the hosting media outlets.

Our experiment focuses on *China Watch*, arguably the largest effort at political native advertising in the world. The survey experiment, which uses real *China Watch* articles as treatments, was conducted in March 2018. Respondents were recruited from Amazon's Mechanical Turk (MTurk) and were directed to an external Qualtrics survey where the experiment took place. A common concern with recruiting a convenience sample from a crowd-sourcing website such as MTurk is that the respondents may be unrepresentative of the population of interest, leading to low external validity (McDermott, 2011). However, recent studies show that MTurk samples are more representative of the overall population than different types of in-person convenience samples (Paolacci et al., 2011; Berinsky et al., 2012). Importantly for external validity, several studies in political science, law, and psychology have also shown that the magnitude of average treatment effects estimated from MTurk samples is similar to the size of effects estimated from nationally representative samples (Berinsky et al., 2012; Clifford et al., 2015; Firth et al., 2017).

Experimental design

To closely resemble the real experience of reading a *China Watch* article, respondents are provided with a link to the hosting media outlet's website to read *one* of two real articles produced by *China Daily*. Some respondents read an article that appeared on the *China Watch* page of the *Washington Post*, while some read an article that appeared on the *China Watch* page of *The Telegraph*. In effect, these particular respondents received a political native advertisement. A third set of respondents read one or the other of these two articles but as it appeared on *China Daily*'s own webpage. In effect, this third set of respondents did not receive a political native advertisement. All respondents received the same message before reading the treatment article: "In the next section, we will ask you to read a short news article. After reading the article, you will be asked to answer several related questions. Please click this link to read the article." We then display the actual URL to the treatment article.[3]

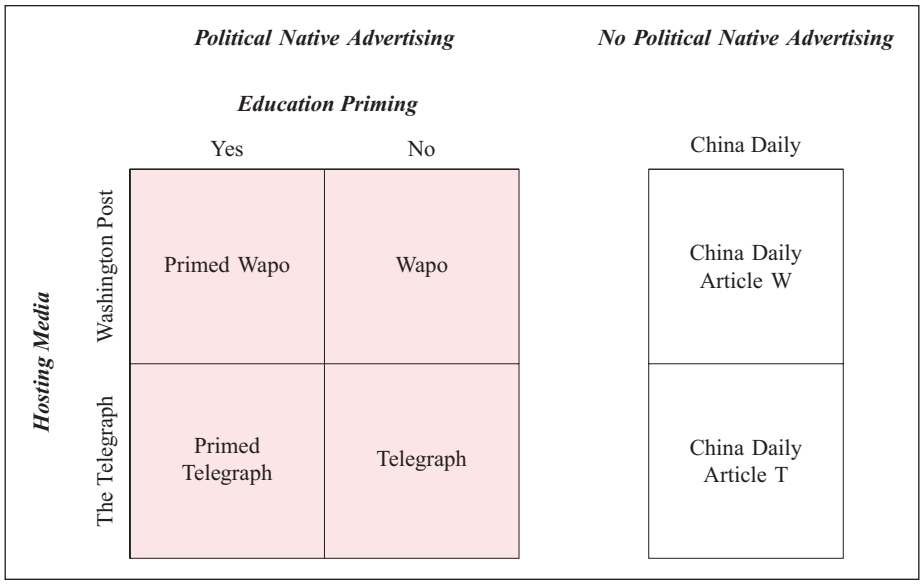

Figure 1. Full experimental design.
Note: Respondents received one of two similar articles related to China's plan to continue with market reforms. These articles appeared in the *Washington Post* or *The Telegraph* (Political Native Advertising) or in *China Daily* (No Political Native Advertising). Half of the respondents who read a political native advertisement were randomly selected to receive an *Education Priming* treatment in which they were told about the practice of political native advertising.

Although it would be ideal for the content of the two *China Watch* articles to be identical, this is not possible. *China Daily* deliberately avoids displaying the exact same articles in different hosting media outlets across countries. We therefore chose similar articles that addressed the same topic: China's plan to continue with market reforms. We chose to use articles that focused on economic news because this is the main type of news story found in the *China Watch* sections of both the *Washington Post* and *The Telegraph*. Importantly, respondent perceptions of the articles they received did not differ across the two *China Watch* articles (see Online Appendix A).

To examine whether increased awareness of political native advertising would increase the likelihood that respondents can detect the true source of the *China Watch* articles, we also include an *Education Priming* treatment related to political native advertising. Specifically, half of the respondents who receive *China Watch* articles are also randomly selected to receive information about the practice of political native advertising. This *Education Priming* treatment comes before the respondents receive their *China Watch* article. The exact wording of the *Education Priming* treatment is

> Foreign countries, such as China and Russia, have started to launch international propaganda campaigns using a technique called native advertising. For example, they buy space on western mainstream media outlets to publish government-sponsored content produced by government-owned media, such as *China Daily*.

A graphical overview of the full experimental design is shown in Figure 1. A total of 660 respondents were randomly assigned among each of the six groups. Two groups of respondents received *Education Priming* before reading the *China Watch* articles. Two groups of respondents read the same *China Watch* articles on China Daily.

Detecting political native advertising. To examine whether respondents are able to detect political native advertising as non-editorial content, respondents who receive one of the *China Watch* articles are asked to identify the source of the article immediately after reading it (left panel of Figure 1). They are given four options: *New York Times, China Daily, Washington Post*, and *The Telegraph*. Studies on commercial native advertising find readers are more likely to recognize native advertising when the disclosure is more explicit, higher in prominence, and positioned at the bottom (Amazeen and Wojdynski, 2018; Amazeen and Muddiman, 2018; Wojdynski and Evans, 2016). Due to different disclosure rules, there is considerable variation in both the clarity and style of disclosure regarding the *China Watch* pages across the *Washington Post* and *The Telegraph*. Screenshots of the *China Watch* pages on the two hosting media outlets are shown in Figure 2. In *The Telegraph*, the *China Watch* page is located under *The Telegraph*'s World News section and closely resembles its own editorial content. Above the *China Watch* articles, there is a statement in small font that reads, "This content is produced and published by *China Daily*, People's Republic of China, which takes sole responsibility for its content." Significantly, there is no indication that the articles are paid supplements from a foreign government. The disclosure in the *Washington Post* is clearer and more prominent than that found in *The Telegraph*.

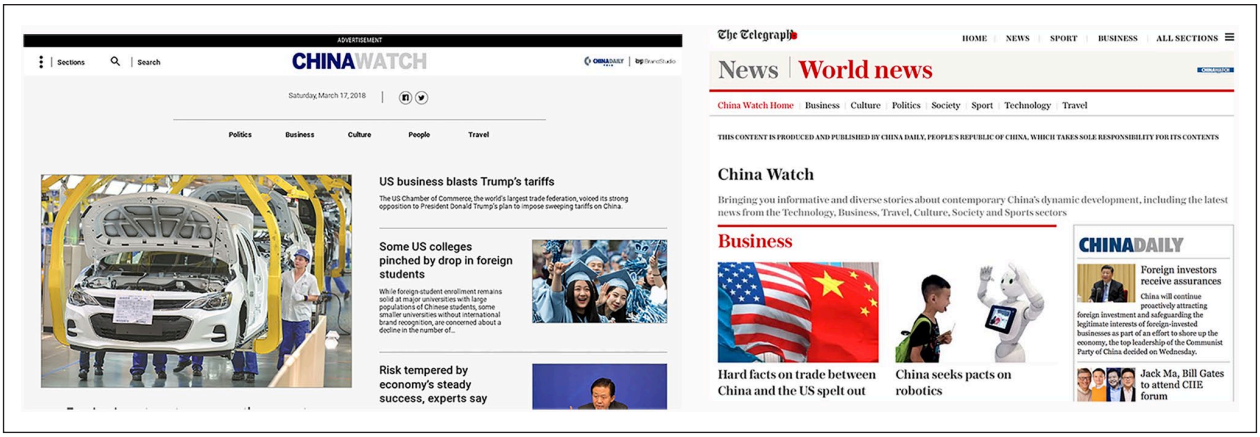

Figure 2. *China Watch* in the *Washington Post* (left) and *The Telegraph* (right).

More importantly, although the web address for the *China Watch* page clearly links it with the *Washington Post*, the material itself looks less like the standard editorial content produced by the *Washington Post*. The *Washington Post* also uses the term "advertisement" at the top of the page and additional information at the bottom of the page states that: "This content is paid for and provided by an advertiser, and the site is managed by WP BrandStudio. The Washington Post newsroom and WP BrandStudio were not involved in the creation of this content." Given the variation in the clarity of disclosure, we expect that respondents who receive a *China Watch* article in the *Washington Post* will be more likely to detect the political native advertising than the respondents who receive a *China Watch* article in *The Telegraph*.

Those who receive the *Education Priming* treatment should be more likely to detect the political native advertisement irrespective of where they see the *China Watch* article. This is because they are primed to look for native advertising, especially as it relates to China. Whether this is in fact the case is important as it speaks to the possibility that information campaigns can immunize readers against the deceptive nature of political native advertising.

Reputational costs of political native advertising. Given that respondents may be unable to detect political native advertising on their own, we inform the respondents that the *China Watch* story they have just read is a paid supplement from the Chinese government-controlled *China Daily* to examine whether there are costs to the hosting media's reputation for publishing political native advertisements. To measure possible reputational costs, we compare the level of respondent trust in the hosting media site *before* they receive their *China Watch* article with their level of trust in the hosting media site *after* they are informed that the article was a paid supplement from a foreign government. A respondent's trust in the hosting media site is measured on a 1–6 scale, where 1 indicates that the respondent has no trust in the hosting media site at all and 6 indicates that the respondent has a great deal of trust in the hosting media site.

While we expect that all respondents will lower their reported trust in the hosting media site when they are told that it is accepting political native advertisements from a foreign government, it is possible that this negative effect will be larger for those respondents who were unable to detect the political native advertising for themselves. These individuals learn that they have been deceived by the hosting media outlet and the foreign government advertiser into thinking that the political advertisement is editorial content by the hosting media. This additional sense of deception may cause their drop in trust in the hosting media outlet to be particularly marked.

Persuasiveness of political native advertising. To examine the persuasiveness of political native advertising, we compare how convincing the respondents find the two *China Watch* articles in the *Washington Post* and *The Telegraph* with how convincing they find the exact same articles when published in *China Daily*. Respondents are asked to rate how convincing their article is on a scale of 1–5, where 1 indicates that the article is not at all convincing and 5 indicates that the article is very convincing. From the perspective of a political actor, the primary appeal of political native advertising is that their message is more convincing if it is camouflaged as standard editorial content on an independent, and thus more credible, hosting media site. Note, though, that the increased persuasiveness of political native advertising ultimately relies on successful deception. If respondents are able to identify the political native advertisement and hence the true source of the news story, then there is no reason to believe that they will find the *China Watch* articles any more (or less) convincing than the same articles in *China Daily*. Thus, political native advertising should only increase the persuasiveness of a news story

among those respondents who fail to detect the political native advertisement.

Experiment results

To make valid causal inferences, respondents must be assigned randomly to different treatment conditions. In other words, there should not be any factors that influence the treatment assignments (Trochim and Donnelly, 2006). Although all of our treatments were randomly assigned, it is possible that the respondents in the various treatment groups could still differ demographically. To evaluate this, we conducted difference-in-means tests on reported demographics among the six treatment groups. We find that the six treatment groups are mostly balanced in terms of age, gender, education, and income. There is some evidence that the respondents in the Education Telegraph and Telegraph groups are slightly younger on average than the respondents in the Primed Wapo group.[4] As a result, we controlled for age in relevant analyses.

Detecting political native advertising

To examine the factors that influence the probability that respondents are able to identify the true source of the *China Watch* story, we estimated a logit regression model in which the dependent variable, *Correct Source*, was coded 1 if the respondent identified the source of the news story as *China Daily* and 0 otherwise. The results from two slightly different model specifications are shown in Table 1.[5]

In Model 1, we find that the effect of the hosting media site, and implicitly the clarity of disclosure, is substantively large and statistically significant. The predicted probability that a respondent is able to identify the true source of the *China Watch* article increases by 0.569 [0.503, 0.631] if a respondent sees the article in the *Washington Post* as opposed to *The Telegraph*.[6] Two-tailed 90% confidence intervals are shown in square brackets. In fact, 72% of respondents who read the *China Watch* article in the *Washington Post* were able to identify the true source of the article, while only 14% of respondents who read the *China Watch* article in *The Telegraph* were able to do so. Although statistically significant, the substantive effect of the *Education Priming* treatment is smaller. To be precise, the predicted probability that a respondent is able to identify the true source of the *China Watch* article increases by 0.075 [0.003, 0.151] when a respondent receives the *Education Priming* treatment.[7]

It is possible that our respondents might be more familiar with the *Washington Post* than *The Telegraph* given that they are US residents, although they indicate a similarly high level of trust toward *The Telegraph* in the survey. The potential extra experience with the *Washington Post* might help the respondents better recognize the stylistic difference between the *China Watch* page and the editorial page in the *Washington Post*. We might expect respondents with more media experience and higher media literacy to be more likely to recognize political native advertising. Therefore, in Model 2 we examine whether a respondent's general level of media experience, media literacy, media trust, and level of education affect the probability that they are able to identify the true source of the *China Watch* article. We employ standard media and education measures from the existing literature. *Media Experience* is based on five dimensions (news consumption, experience, expertise, familiarity, and access), each of which is measured on a five-point scale. *Media Experience* is calculated as the sum of the scores for each of these five dimensions; its possible values run from 5 to 25, with higher numbers indicating more media experience (Flanagin and Metzger, 2000). *Media Literacy* captures a respondent's understanding of media ownership, media and politics, media effects, news framing, agenda setting, and the role of journalists. It draws on 14 survey items that are each measured on a five-point scale. *Media Literacy* is calculated as the sum of the scores

Table 1. Identifying the correct source for the *China Watch* news story.

	Dependent Variable: Correct Source (0,1)	
	Model 1	Model 2
Washington Post	2.726***	2.854***
	(0.241)	(0.252)
Education Priming	0.413*	0.468*
	(0.238)	(0.246)
Media Experience		0.028
		(0.041)
Media Literacy		−0.022
		(0.035)
Media Trust		0.049
		(0.048)
Age	0.011	0.012
	(0.010)	(0.011)
Education		0.128
		(0.097)
Income		0.038
		(0.039)
Female		−0.110
		(0.252)
Constant	−2.271***	−3.754*
	(0.473)	(2.045)
Observations	444	439
Log likelihood	−222.462	−213.998
Akaike Inf. Crit.	452.923	447.995

*$p<0.1$; **$p<0.05$; ***$p<0.01$ (two-tailed).
Note: The dependent variable, *Correct Source*, is a dichotomous variable that equals 1 if a respondent correctly identifies the true source of the *China Watch* article as *China Daily*, and 0 otherwise. Only those respondents who received a *China Watch* article are included. Estimates are based on logit regressions. Standard errors are shown in parentheses.

Table 2. Comparing pre- and post-treatment trust in the hosting media outlets.

	Trust in hosting media outlet (1–6)			
	Pre-treatment	Post-treatment	Paired difference (post–pre)	N
Washington Post	4.09 (1.36)	3.29 (1.33)	−0.80***	182
The Telegraph	3.42 (1.36)	2.65 (1.31)	−0.77***	121

*p<0.1; **p<0.05; ***p<0.01.

Note: The pre-treatment and post-treatment columns indicate the mean level of trust in each of the two hosting media sites before the respondents receive their *China Watch* article and after respondents are informed that the article is a paid supplement from a foreign government. Standard deviations are shown in parentheses. The paired difference column indicates the mean of the paired differences for each respondent between their pre- and post-treatment levels of trust. Only the respondents with complete answers to both pre- and post-treatment questions on the trust in hosting media outlets that are used for the paired t-tests are included to calculate the pre-treatment and post-treatment group means.

for each of these items; its possible values run from 14 to 70, with higher numbers indicating greater media literacy (Ashley et al., 2013). *Media Trust* uses seven survey items, each measured on a five-point scale, to measure an individual's level of trust in the general media (Tsfati and Cappella, 2003). *Media Trust* is simply the sum of the scores on the seven survey items; its possible values run from 7 to 35, with higher numbers indicating greater trust in the news media in general. *Education* is measured on a 1–8 scale, with larger numbers indicating higher levels of education. As the results in Model 2 indicate, none of these variables have any significant effect on the probability that a respondent is able to identify the true source of the *China Watch* articles. A likelihood ratio test indicates that the media and education variables are also jointly insignificant. These variables are also insignificant if we remove our two treatment variables, *Washington Post* and *Education Priming*. In sum, there is no evidence that more educated or more media-savvy individuals are any better at identifying political native advertising than other individuals.

Reputational costs of political native advertising. In Table 2, we compare the mean level of trust in each of the two hosting media sites before the respondents receive their *China Watch* article (Pre-Treatment) and after the respondents are informed that the article is a paid supplement from a foreign government (Post-Treatment). The third column, Paired Difference, indicates whether the mean of the paired differences in the pre-treatment and post-treatment levels of trust for each respondent are statistically different.[8] As expected, trust in both the *Washington Post* and *The Telegraph* declines in a statistically significant way when respondents are told that the *China Watch* article comes from the Chinese government-controlled *China Daily* and not the hosting media outlet. The magnitude of the decline in trust is almost identical for both hosting media outlets — the average decline in trust is 0.80 points for those respondents who received the *China Watch* article in the *Washington Post* and 0.77 points for those respondents who received the article in *The Telegraph*. These differences are substantively meaningful as they equate to a 27.4% decline in trust for the *Washington Post* and a 29.2% decline in trust for *The Telegraph*. The fact that the decline in trust is similar across the two hosting media outlets is interesting as it suggests that the better disclosure provided by the *Washington Post* does not immunize it against potential reputational costs if the political native advertising is detected, which contradicts findings in commercial native advertising that readers have less unfavorable view toward the hosting media outlet when it is more transparent about native advertising (Amazeen and Wojdynski, 2019). It is possible that news consumers are much less tolerant when it comes to political native advertising, especially when the advertisers are foreign governments, given the political implications that independent media outlets can be co-opted by foreign governments to deliver propaganda.

Recall that we expect respondents who are unable to detect political native advertising on their own to exhibit a particularly marked reduction in their trust toward the hosting media outlet after they are told that the *China Watch* article is a paid advertisement from the Chinese government due to the additional perceived deception. In line with our expectation, the results reported in Table 3 indicate that the decline in trust toward the hosting media outlet is in fact larger among the respondents who were unable to detect the political native advertising for themselves than among those respondents who were able to detect it. This is true regardless of whether the respondents saw the *China Watch* article in the *Washington Post* or *The Telegraph*. In robustness checks reported in Online Appendix A, we find that this effect of self-detection is also statistically significant after controlling for demographic and media-related variables. Together the results show that hosting media outlets that accept political native advertising can expect to suffer a reputational cost if the political native advertising is revealed, especially if readers feel that they have been deceived.[9]

Persuasiveness of political native advertising

To examine the persuasiveness of political native advertising, in Table 4 we compare how persuasive the respondents find the two *China Watch* articles based on whether they

Table 3. Change in trust in the hosting media outlets.

		Trust in hosting media outlet (1–6)		
		Pre-treatment	Post-treatment	Paired difference (post–pre)
Washington Post	Self-detected	4.09 (1.39)	3.33 (1.36)	−0.76***
	Not self-detected	4.08 (1.25)	3.15 (1.25)	−0.93***
The Telegraph	Self-detected	3.36 (1.47)	2.73 (1.16)	−0.64
	Not self-detected	3.43 (1.34)	2.64 (1.35)	−0.80***

*p<0.1; **p<0.05; ***p<0.01.

Note: The pre-treatment and post-treatment columns indicate the mean level of trust in each of the two hosting media sites before the respondents receive their *China Watch* article and after respondents are informed that the article is a paid supplement from a foreign government. The two rows for each media outlet distinguish between those respondents who were able to detect the political native advertisement on their own and those respondents who were not. Standard deviations are shown in parentheses. The paired difference column indicates the mean of the paired differences for each respondent between their pre- and post-treatment levels of trust. Only the respondents with complete answers to both pre- and post-treatment questions on the trust in hosting media outlets that are used for the paired t-tests are included to calculate the pre-treatment and post-treatment group means.

Table 4. Persuasiveness, deception, and political native advertising.

	Mean of persuasiveness (1–5)			
	On hosting media		In *China Daily*	Difference
	Identified true source	Not identified true source		
Washington Post China Watch article	3.49 (0.83)		3.47 (0.89)	0.02
		3.77 (0.84)	**3.47 (0.89)**	**0.31***
The Telegraph China Watch article	3.27 (0.96)		3.43 (1.07)	–0.17
		3.87 (0.84)	**3.43 (1.07)**	**0.44*****

*p<0.1; **p<0.05; ***p<0.01.

Note: The first two columns indicate how persuasive the respondents found each of the two *China Watch* articles on average when they read the stories on the hosting media outlet. The first column focuses on those respondents who identified the true source of the *China Watch* article, while the second column focuses on those respondents who failed to identify the true source of the article. Standard deviations are included in parentheses. The third column indicates how persuasive the respondents found each of the two articles when they read the stories in *China Daily*. The final column indicates whether the mean level of persuasiveness was significantly higher when the articles were read on the hosting media outlet instead of *China Daily*. Cells shown in bold refer to those respondents who failed to detect the political native advertisement.

saw the article on one of the hosting media outlets (Political Native Advertising) or in *China Daily* (No Political Native Advertising) and whether they were able to detect the native advertising on their own. Note that we exclude the respondents who received *Education Priming* in this analysis. Consistent with our expectations, the articles that appear on the hosting media site are considered significantly more persuasive than the same article in *China Daily* when the respondents fail to identify the true source of the article. This is true both for those respondents who saw the article in the *Washington Post* and those who saw the article in *The Telegraph*: the *China Watch* article that is perceived as from the *Washington Post* is, on average, 0.31 points more persuasive (9% increase) than the same article in *China Daily*. The *China Watch* article that is perceived as from *The Telegraph* is, on average, 0.44 points more persuasive (13% increase) than the same article in *China Daily*. This is exactly what we expected and highlights the fact that the effectiveness of political native advertising relies on successful deception.[10]

Conclusion

In this article, we study an overlooked yet important form of propaganda, political native advertising, in which political actors buy space on independent media sites to publish political advertisements that mimic the standard editorial content found on the hosting media site. As one of the first empirical studies on political native advertising, we examine some basic but important key features of this type of propaganda. Specifically, we examine the deceptiveness and persuasiveness of political native advertisements and the reputational costs associated with political native advertising to the hosting media outlets. Using an online survey experiment with real political native advertisements, we find that unless the political native advertising is clearly disclosed using visual separation as in the *Washington Post*, respondents are unlikely to detect the true source of political native advertising, irrespective of their level of education and experience as a news consumer. The deceptive nature of political native advertising

makes it an effective tool for political actors to influence the citizenry. As we show in our results, the same message is perceived as much more persuasive when it comes in the form of native advertisement and is perceived as coming from the independent hosting media outlet. Note that the respondents are recruited from MTurk, which are previously shown to be better educated and more media savvy than the overall population in the US (Berinsky et al., 2012), and therefore the effects found in the experiment are likely to be even stronger in the population: citizens, on average, are likely to have an even harder time detecting political native advertising than the respondents in the experiment. The increased salience of foreign disinformation campaigns after the 2016 election in the USA is also likely to make citizens more alert regarding political native advertising and likely to punish the hosting media outlets more than citizens in other targeted countries. However, with the increased popular and scholarly concern over disinformation and misinformation campaigns around the world (e.g. Brady, 2015, 2018; Tucker et al., 2017), we are likely to observe similar responses to political native advertising from citizens outside of the USA.

Although our findings are troubling in that citizens are often unable to detect political native advertising unless it is clearly disclosed, and are likely to be influenced by such propaganda, our analysis also suggests methods to deter such propaganda. Much of the effectiveness of political native advertising depends on the deceptiveness of the advertising. It suggests that decreasing the deceptiveness of the advertising and/or increasing the reader's ability to detect the political advertising are likely to reduce the incentives that political actors have to engage in these sorts of propaganda activities and deter independent media outlets from cooperating with political actors. However, there seem to be no significant reputational gains to the hosting media outlet to disclose political native advertising. As a result, disclosure regulations are needed to provide media outlets with additional motivations to disclose political native advertising. In addition, although less effective than clear disclosure, informational campaigns that publicize the presence and use of political native advertising can also increase the reader's ability to recognize political native advertising.

Our results also have important regulatory implications for native advertising in general. Although studies on commercial native advertising find certain types of disclosure to be more effective in increasing the respondents' advertisement recognition, in general only around 10% of the respondents were able to detect native advertising under the better disclosure conditions (Amazeen and Wojdynski, 2019; Wojdynski and Evans, 2016). Our finding that respondents were much more likely to detect the true source of the *China Watch* article in the *Washington Post* suggests that increasing the degree of visual separation between the native advertisement and editorial content might be a better method than simply using text disclosures to alert readers of native advertising.

Acknowledgements

We thank Amber Boydstun, Matt Golder, Sona Golder, Arthur Lupia, Benjamin Radford, David Szakonyi, Rory Truex, and the anonymous reviewers, as well as audiences at the 2019 Annual Meeting of the Southern Political Science Association and the 2019 Annual Meeting of the American Political Science Association for their helpful comments on this paper.

Declaration of conflicting interests

The author(s) declared no potential conflicts of interest with respect to the research, authorship, and/or publication of this article.

Ethical Approval

The online survey experiment described in this paper received The Institutional Review Board (IRB) approval (Study #00008806) at Pennsylvania State University.

Funding

The author(s) received no financial support for the research, authorship, and/or publication of this article.

ORCID iD

Yaoyao Dai https://orcid.org/0000-0003-3035-331X

Supplemental materials

1. The supplemental files are available at http://journals.sagepub. com/doi/suppl/10.1177/2053168020935250
2. The replication files are available at https://dataverse.harvard.edu/dataset.xhtml?persistentId=doi:10.7910/DVN/GBLKVP.

Notes

1. Legacy media refer to older and more traditional media outlets such as newspapers, television, and radio, in which the audience does not actively "interact" with the media content.
2. The hosted website and column, russianow.washingtonpost.com, disappeared in 2015. For information and reports on *Russia Now*, see Barton (2015) and the Washington Free Beacon (2014).
3. For example, the URL for the Washington Post's *China Watch* article is http://chinawatch.washingtonpost.com/2018/02/capital-market-to-open-further/. This URL no longer works, since the *Washington Post* removed the *China Watch* page in the summer of 2018, while continuing to deliver the print version of *China Watch* with the printed *Washington Post.*
4. Detailed results are included in Online Appendix A section 1.
5. Note that only the 444 respondents who are in the four treatment groups with *China Watch* articles on the hosting media outlets are relevant and hence included in this analysis.
6. This change in predicted probability is calculated for a respondent who is 41 years old (sample mean) and who did not receive the *Education Priming* treatment.

7. This change in predicted probability is calculated for a respondent who is 41 years old and who saw the *China Watch* article on the *Washington Post* site.
8. The numbers in the Paired Difference column reflect the average difference in the *paired* post-treatment and pre-treatment levels of trust from *each* respondent which is different from the difference in the two group means of the post-treatment levels of trust and pre-treatment levels of trust from *all* respondents. A paired difference-in-means test is appropriate for our within-subject before-and-after experimental design and is generally more powerful than an unpaired test because it reduces intersubject variability (Larsen et al., 1986).
9. Additional robustness checks controlling for demographic and media related variables are included in Online Appendix A, Section 4.
10. Additional robustness checks controlling for demographic and media related variables are included in Online Appendix A, Section 5.

Carnegie Corporation of New York Grant

This publication was made possible (in part) by a grant from the Carnegie Corporation of New York. The statements made and views expressed are solely the responsibility of the author.

References

Amazeen MA and Muddiman AR (2018) Saving media or trading on trust? The effects of native advertising on audience perceptions of legacy and online news publishers. *Digital Journalism* 6(2): 176–195.

Amazeen MA and Wojdynski BW (2018) The effects of disclosure format on native advertising recognition and audience perceptions of legacy and online news publishers. *Journalism*: 1464884918754829.

Amazeen MA and Wojdynski BW (2019) Reducing native advertising deception: Revisiting the antecedents and consequences of persuasion knowledge in digital news contexts. *Mass Communication and Society* 22(2): 222–247.

Archer AM and Clinton J (2018) Changing owners, changing content: Does who owns the news matter for the news? *Political Communication* 35(3): 353–370.

Ashley S, Maksl A and Craft S (2013) Developing a news media literacy scale. *Journalism & Mass Communication Educator* 68(1): 7–21.

Barton E (2015) China, Russia pay *Washington Post* to publish their propaganda. *The Daily Caller*, October 29.

Batsell J (2017) 4 steps to bring ethical clarity to native advertising. NiemanReports. Harvard University: *Nieman Foundation*.

Berinsky AJ, Huber GA and Lenz GS (2012) Evaluating online labor markets for experimental research: Amazon.com's mechanical turk. *Political Analysis* 20(3): 351–368.

Brady AM (2015) China's foreign propaganda machine. *Journal of Democracy* 26(4): 51–59.

Brady AM (2018) China in Xi's "new era": New Zealand and the CCP's "magic weapons." *Journal of Democracy* 29(2): 68–75.

Carlson M (2015) When news sites go native: Redefining the advertising–editorial divide in response to native advertising. *Journalism* 16(7): 849–865.

China Daily (2018) *China Daily*'s global development. Available at: http://www.chinadaily.com.cn/static_e/global.html (accessed 12 June 2020).

Clifford S, Jewell RM and Waggoner PD (2015) Are samples drawn from mechanical turk valid for research on political ideology? *Research and Politics* 4(2): 1–9.

Cook S (2017) Chinese government influence on the U.S. media landscape. Testimony before the U.S.-China Economic and Security Review Commission Hearing on China's Information Controls, Global Media Influence, and Cyber Warfare Strategy, May 4.

Crabtree C, Kern HL and Siegel DA (2018) Cults of personality, preference falsification, and the dictator's dilemma. Unpublished manuscript, University of Michigan.

Di Tella R and Franceschelli I (2011) Government advertising and media coverage of corruption scandals. *American Economic Journal: Applied Economics* 3(4): 119–151.

Dykhne I (2017) Persuasive or deceptive: native advertising in political campaigns. *Southern California Law Review* 91: 339.

Edmonds R (2017) Native advertising grows up fast, shedding its rogue image. *Poynter*. Available at: https://www.poynter.org/business-work/2017/native-advertising-grows-up-fast-shedding-its-rogue-image/ (accessed 12 June 2020).

Einstein M (2016) *Black Ops Advertising: Native Ads, Content Marketing and the Covert World of the Digital Sell*. New York: OR Books.

Fallows J (2011) Official Chinese propaganda: Now online from the WaPo! *The Atlantic*, February 3.

Firth K, Hoffman DA and Wilkinson-Ryan T (2017) Law and psychology grows up, goes online and replicates. Public Law and Legal Theory Research Paper Series # 17–38, University of Pennsylvania Law School.

Flanagin A and Metzger M (2000) Perceptions of internet information credibility. *Journalism & Mass Communication Quarterly* 77(3): 515–540.

Gehlbach S and Sonin K (2014) Government control of the media. *Journal of Public Economics* 118: 163–171.

Howe P and Teufel B (2014) Native advertising and digital natives: The effects of age and advertisement format on news website credibility judgments. *ISOJ Journal* 4(1): 78–90.

Huang H (2015) Propaganda as signaling. *Comparative Politics* 47(4): 419–444.

Huang H (2018) The pathology of hard propaganda. *Journal of Politics* 80(3): 1034–1038.

Iversen MH and Knudsen E (2017) When politicians go native: The consequences of political native advertising for citizens trust in news. *Journalism*: 1464884916688289.

Jamieson KH, et al. (2000) *Everything You Think You Know about Politics—and Why You're Wrong*. New York: Basic Books.

Larsen RJ, Marx ML, et al. (1986) *An Introduction to Mathematical Statistics and Its Applications*, vol. 2. Englewood Cliffs, NJ: Prentice-Hall.

Little A (2018) Fake news, propaganda, and lies can be pervasive even if they aren't persuasive. In: Golder M and Golder SN (eds.) *Fake News and the Politics of Misinformation, Comparative Politics Newsletter* 28(2), 49–54.

Little AT (2017) Propaganda and credulity. *Games and Economic Behavior* 102: 224–232.

Marquez X (2018) Two models of political leader cults: Propaganda and ritual. *Politics, Religion, and Ideology* 19(3): 265–284.

McDermott R (2011) New directions for experimental work in international relations. *International Studies Quarterly* 55(2): 503–520.

Mullin B (2017) Vox Media pitches signature "Explainer" format to advertisers. 13 October. *Wall Street Journal*. Available at: https://www.wsj.com/articles/vox-media-pitches-signature-explainer-format-to-advertisers-1507892401 (accessed 12 June 2020).

Murtha J and Gourarie C (2016) Do BuzzFeeds native political ads cross a line? *Columbia Journalism Review*. Available at: https://www.cjr.org/analysis/a_major_player_in_the.php (accessed 12 June 2020).

Paolacci G, Chandler J and Ipeirotis PG (2011) Running experiments on Amazon mechanical turk. *Judgment and Decision Making* 5(5): 411–419.

Stockmann D (2013) *Media Commercialization and Authoritarian Rule in China*. Cambridge: Cambridge University Press.

Trochim WM and Donnelly JP (2006) *Research Methods Knowledge Base*. Cincinnati, OH: Atomic Dog.

Tsfati Y and Cappella JN (2003) Do people watch what they do not trust? Exploring the association between news media skepticism and exposure. *Communication Research* 30(5): 504–529.

Tucker JA, Theocharis Y, Roberts ME and Barberá P (2017) From liberation to turmoil: Social media and democracy. *Journal of democracy* 28(4): 46–59.

Washington Free Beacon (2014) WaPo still publishing Russian propaganda. September 3.

Wojdynski BW and Evans NJ (2016) Going native: Effects of disclosure position and language on the recognition and evaluation of online native advertising. *Journal of Advertising* 45(2): 157–168.